


2022 Fall @ POSTECH

EECE695D: Efficient ML Systems

Semester in Review

1. Linear Models



We first started with “**how to count memory / compute**” of both machine learning inference and training.

- **Compute.** FLOPs and MACs...
(rounding issues of MAC, tradeoff with parallel)
- **Memory.** So-called von Neumann bottleneck
(very basic ideas of how CPU works, idle time)
- **Training cost.** Depends on your optimization algorithm!
(Explicit solution vs. approximate solution)
- **Matmuls.** Modern ML is all about matrix multiplications!
(some parts can be precomputed!)


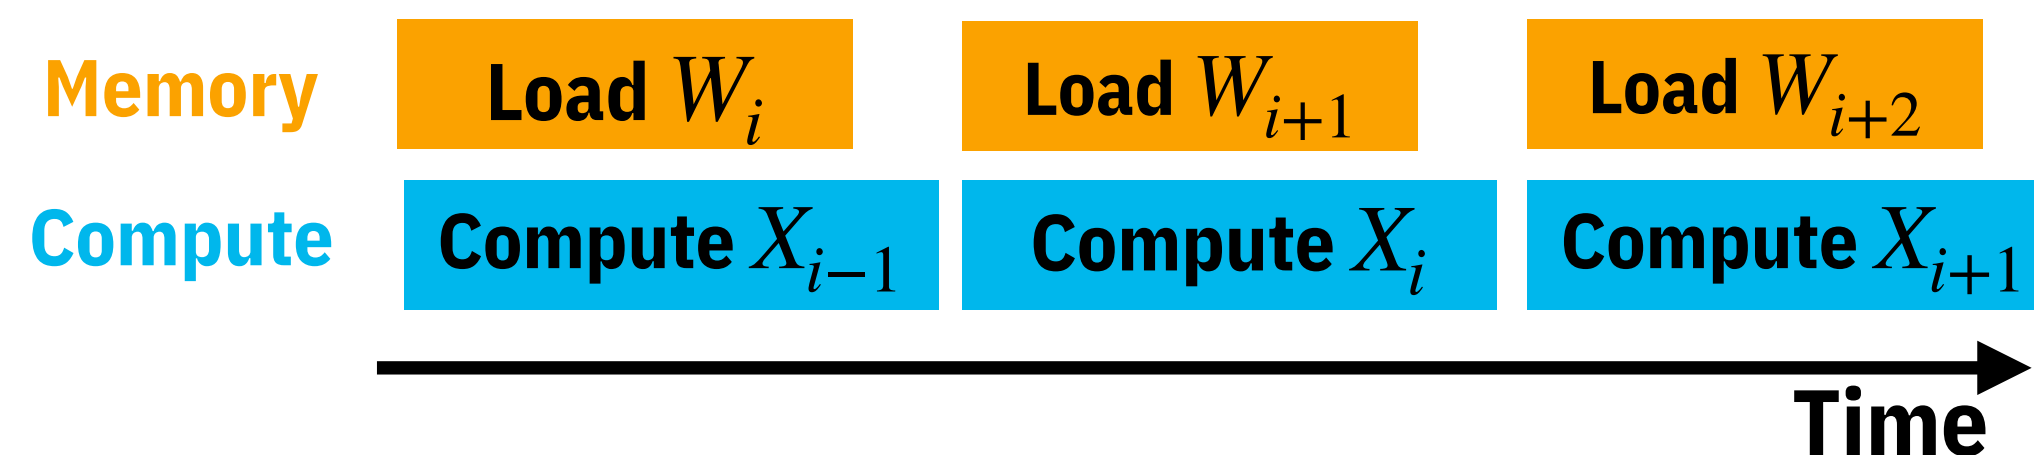
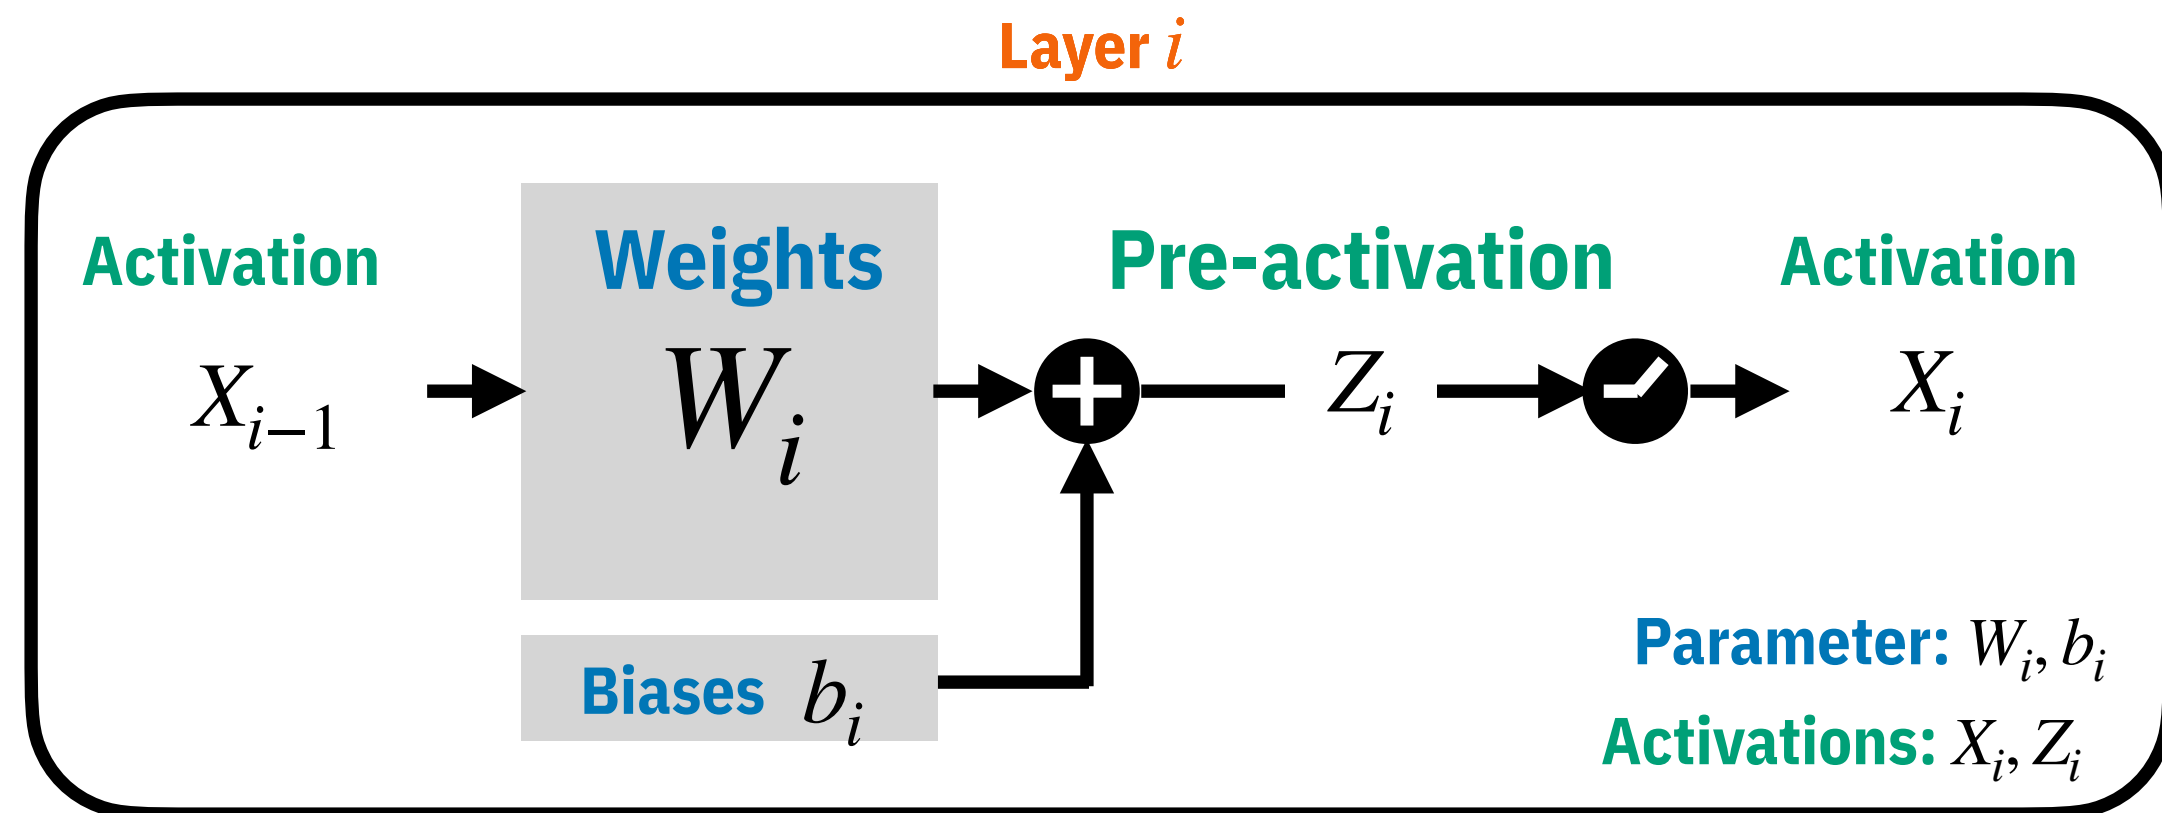
$$\mathbf{w}^{\text{new}} = \mathbf{w} - 2\epsilon \cdot (\mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{X}^T \mathbf{y})$$

Option 1. Precompute $\mathbf{X}^T \mathbf{X}$ and reuse.

Option 2. Compute $\mathbf{X} \mathbf{w}$ first
(better with less #steps).



Note. Only requires computing once!
Can be reused for every iteration.

2. Deep Models




- Repetition of linear models + activation functions
- “Loading all weights” does not work—**Scheduling** of weight loading and computing! (overlapping is a common practice)
- The exact solution does not work—backpropagation!
 $2 * \text{Forward FLOPs} \approx \text{Backward FLOPs}$
(Requires us to keep track of the activations and gradients of other layers... memory checkpointing...)
- Modules with smaller compute/memory footprint, e.g., depthwise convolutions
- Efficiency metrics: Time, #Param, Compute, Energy

3. Bias-Variance Tradeoff...?



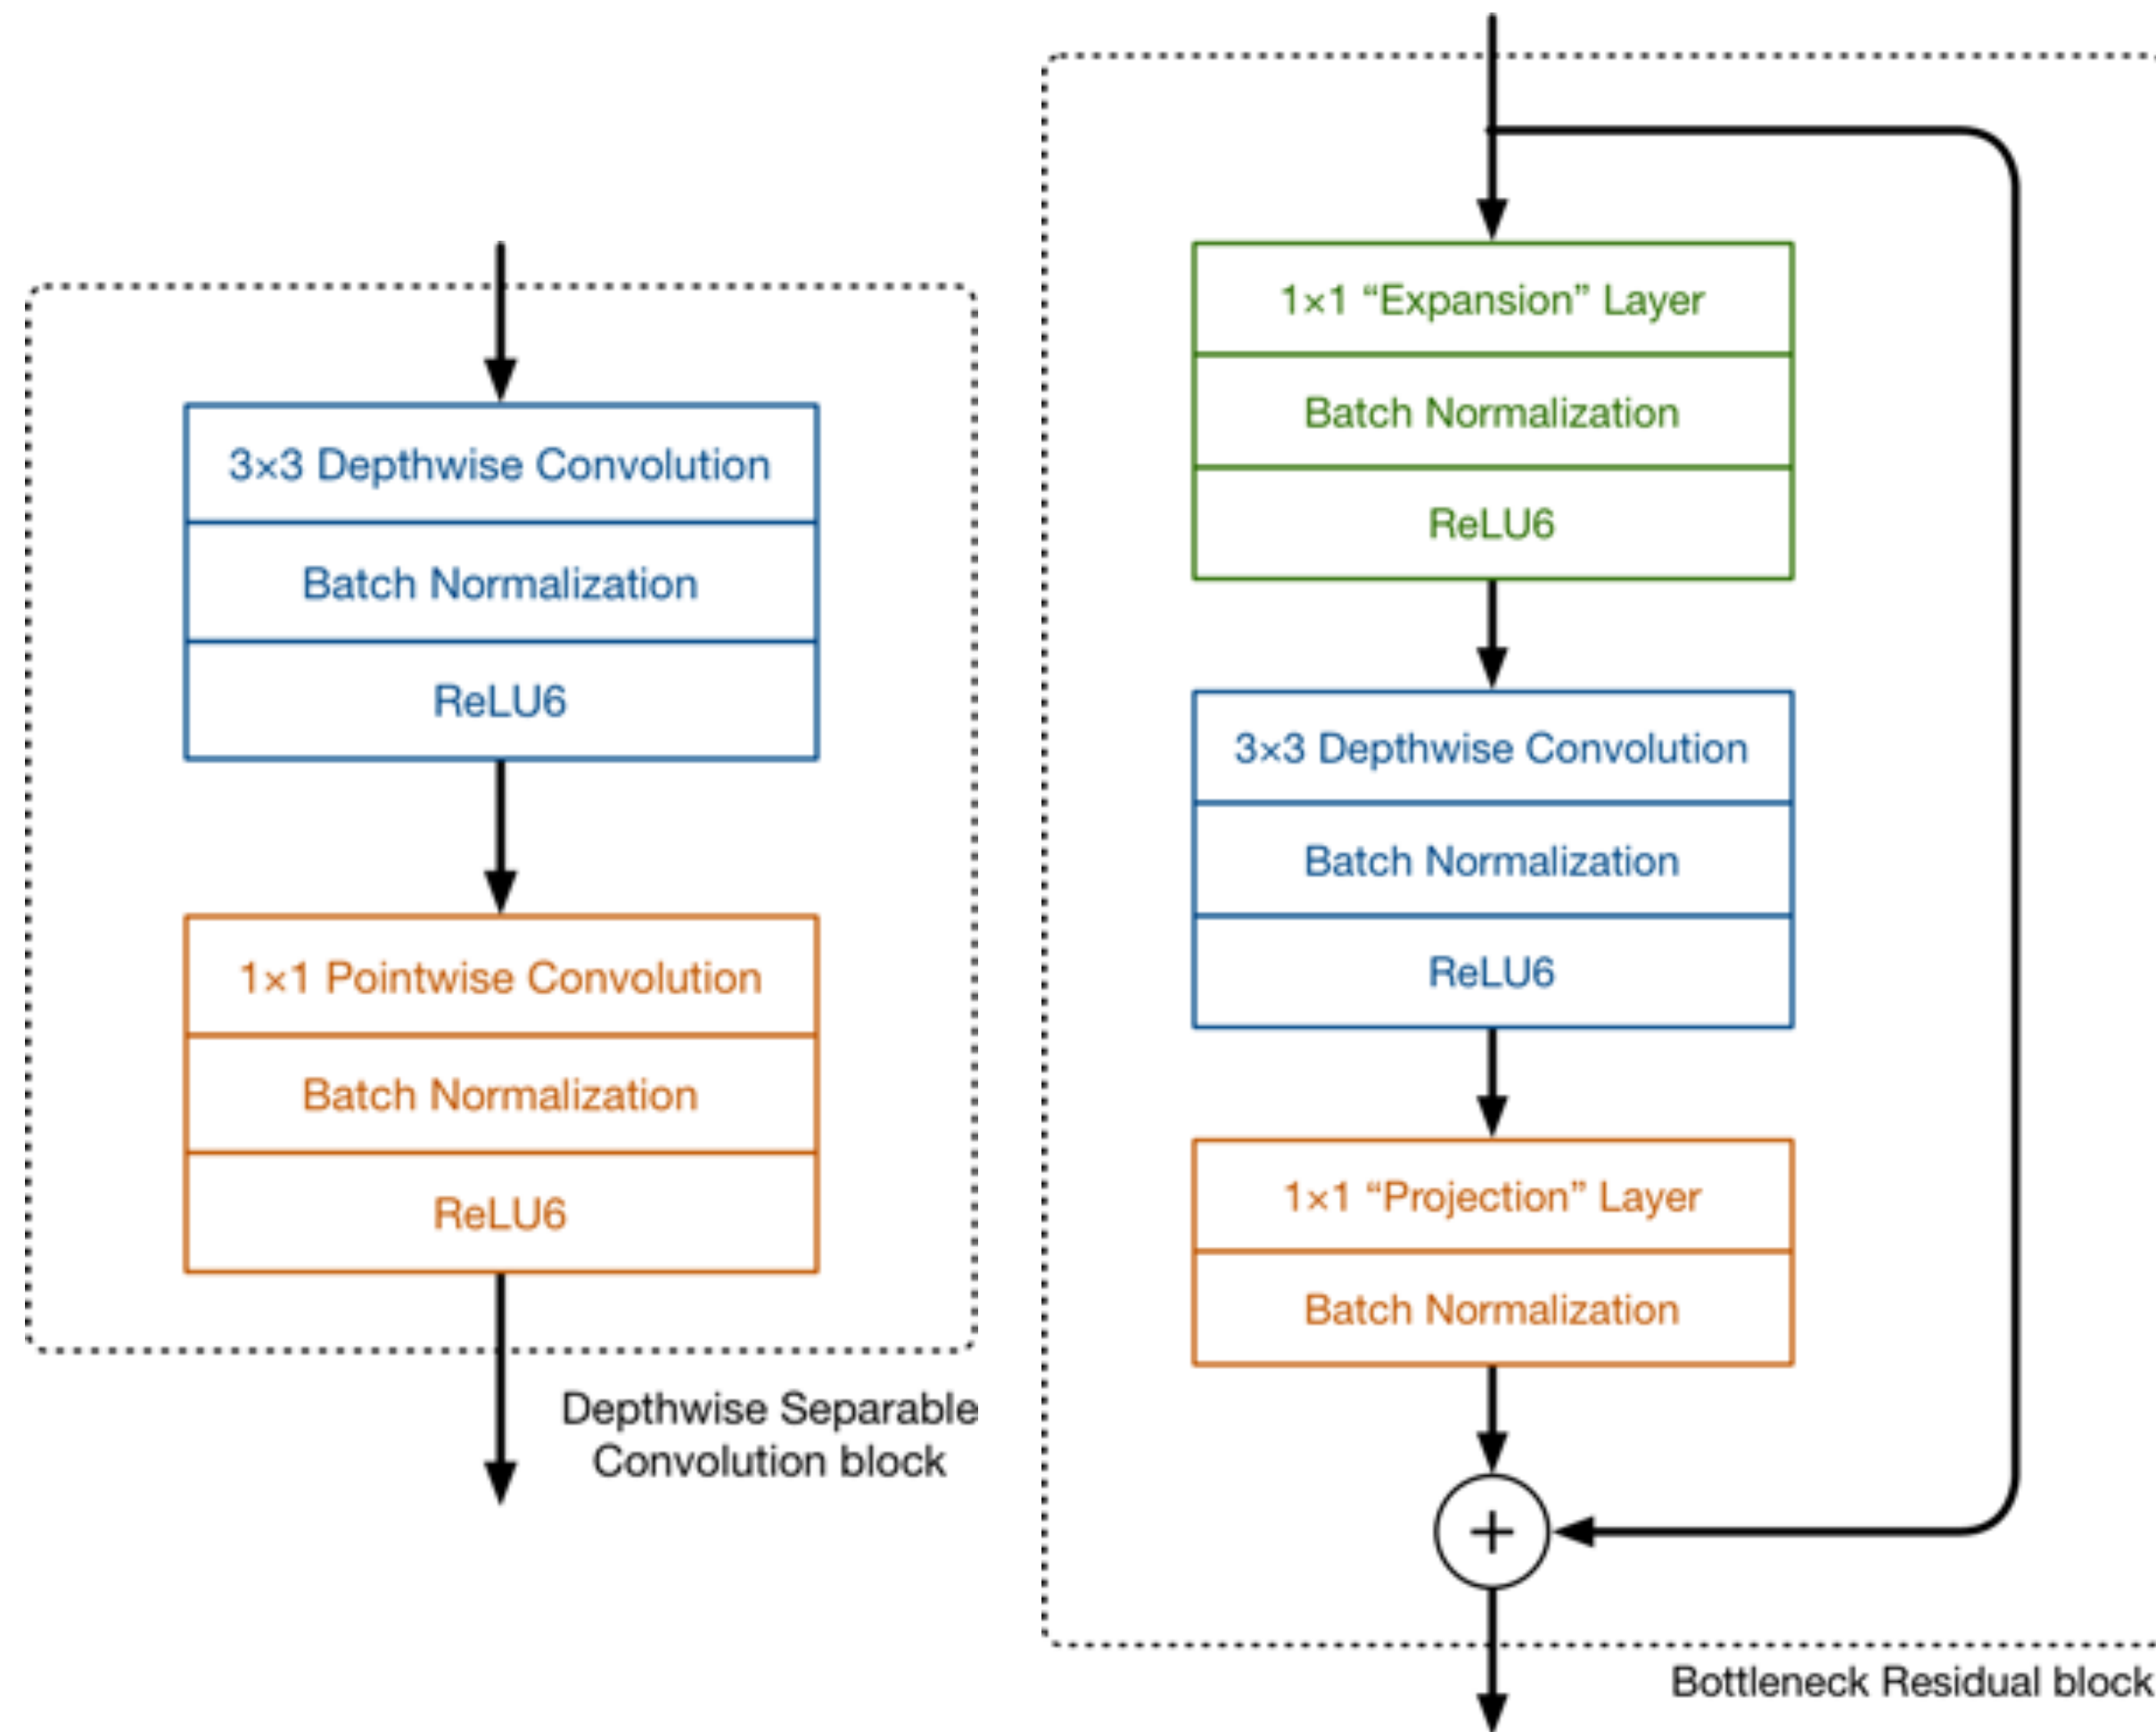
- In the past, people believed that large models (larger than dataset) do not generalize well.
- In these days, people are more sure that larger models generalize better (if they are large enough)
- There are some theories to explain why this happens... (slightly out of our scope though)

4. Neural Scaling Laws



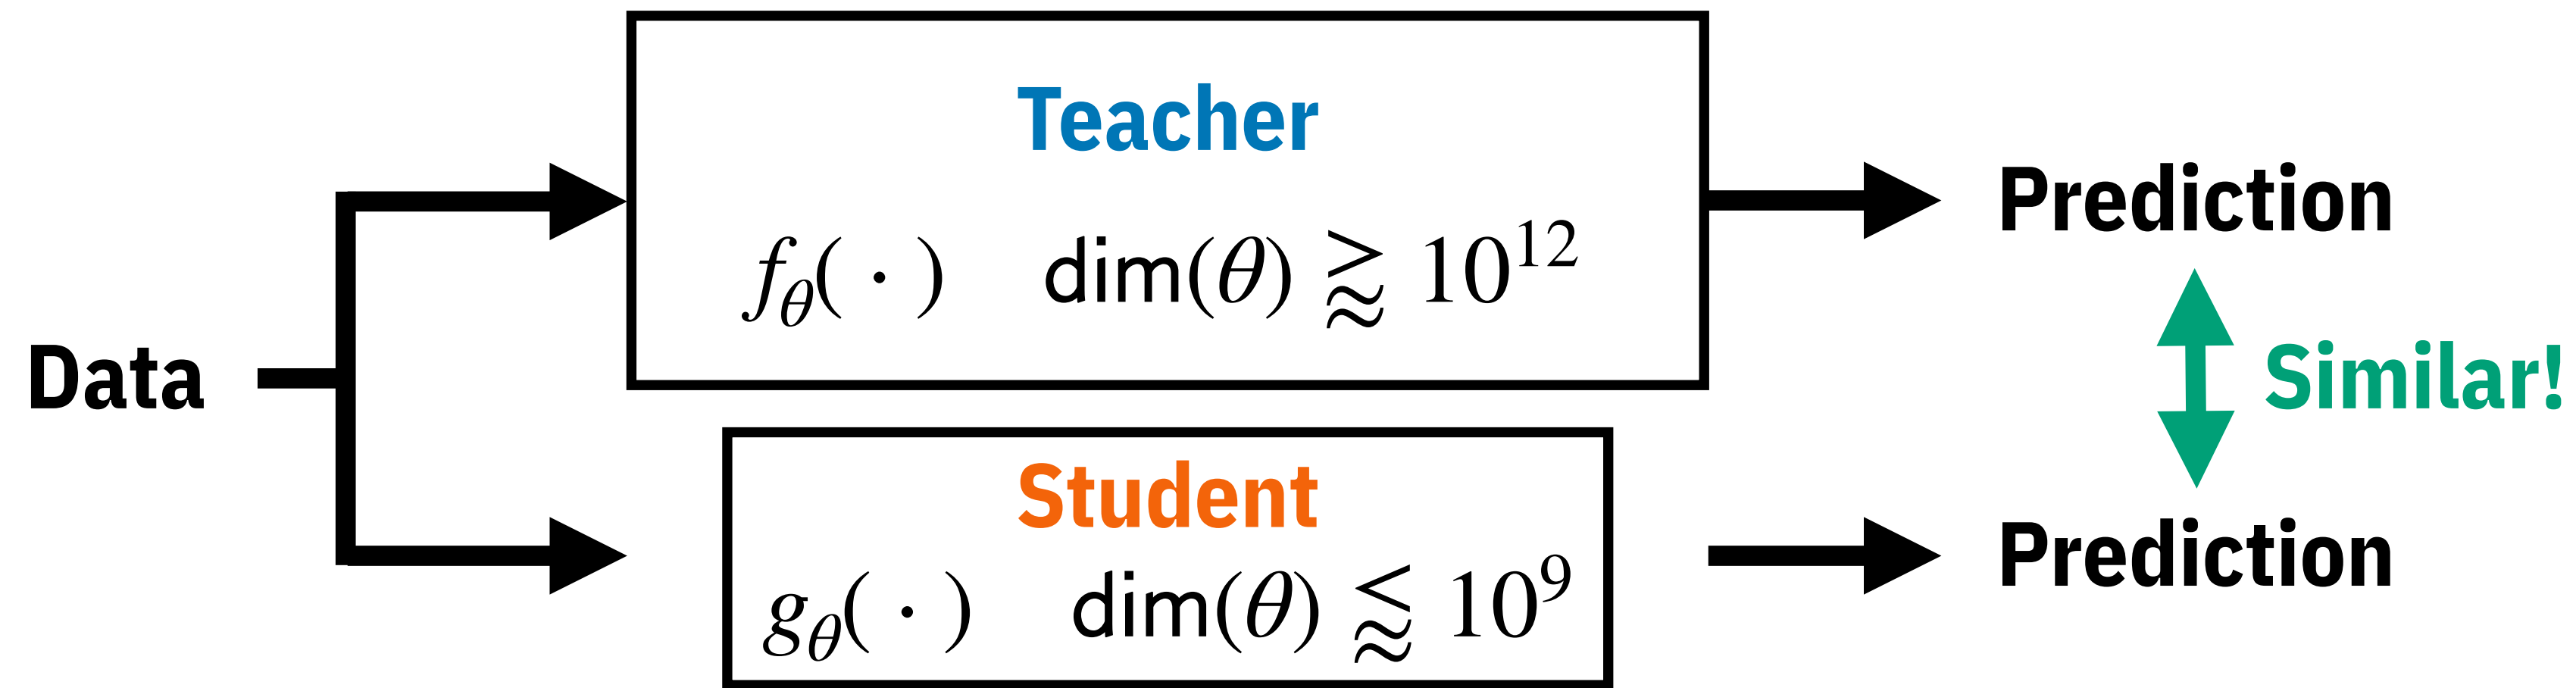
- People find that, neural network scaling follows some “**power law**” (happens for architectures that scale well, e.g., transformers)
- The compute, #parameters, and dataset should grow simultaneously—the performance saturates otherwise.
- Recent works use this model to generate compute-efficient models (e.g., Chinchilla)

5. Model Architectures



- To design compute/memory-efficient models, people used to design new modules.
 - Fire module, with 1x1 conv (e.g., SqueezeNet)
 - Inverted Residual, with depthwise separable conv (e.g., MBNetv2)
 - ReLU6 (e.g., MBNetv1)
- More recently, people used NAS to build new models based on these modules (e.g., MBNetv3)
- The focus is moving toward trainability—via scaling laws (e.g., EfficientNet) via memory-efficiency (e.g., MCUNets)


6. Knowledge Distillation



- Enhance the training of a small model, by training it to **approximate the output of a large model** (or itself)—Benefits in terms of generalization, rather than approximation.
 - Mimic the softmax outputs
 - Mimic the intermediate activations
 - Mimic the relationship between data samples

7. Pruning

$$\begin{bmatrix} a_1 & a_2 & a_3 & a_4 \\ a_5 & a_6 & a_7 & a_8 \\ a_9 & a_{10} & a_{11} & a_{12} \\ a_{13} & a_{14} & a_{15} & a_{16} \end{bmatrix} \longrightarrow \begin{bmatrix} 0 & 0 & \tilde{a}_3 & \tilde{a}_4 \\ \tilde{a}_5 & 0 & \tilde{a}_7 & 0 \\ \tilde{a}_9 & 0 & 0 & \tilde{a}_{12} \\ \tilde{a}_{13} & 0 & \tilde{a}_{15} & \tilde{a}_{16} \end{bmatrix}$$



- Add zeros to the weights of a neural network, so that the associated compute/memory is no longer necessary.
- Key questions are **what to prune, when to prune, how much to prune**.
 - A popular method is to prune by magnitude, gradually by cubic schedule, with heuristic layerwise sparsity (e.g., Gale et al., (2019))
Pruning at initialization, or using Hessian-based criteria is also popular.
- Difficult to exploit the benefit if there is no structure to the zeros—prune whole filters or force patterns.
- Note: There is a recent work that argues “for some tasks, you can’t do pruning” (I don’t agree though)

8. Activation Sparsity

$$\begin{bmatrix} w_1 & w_2 & w_3 \\ w_4 & w_5 & w_6 \end{bmatrix} \begin{bmatrix} x_1 & x_2 \\ x_3 & x_4 \\ x_5 & x_6 \end{bmatrix}$$

Dense

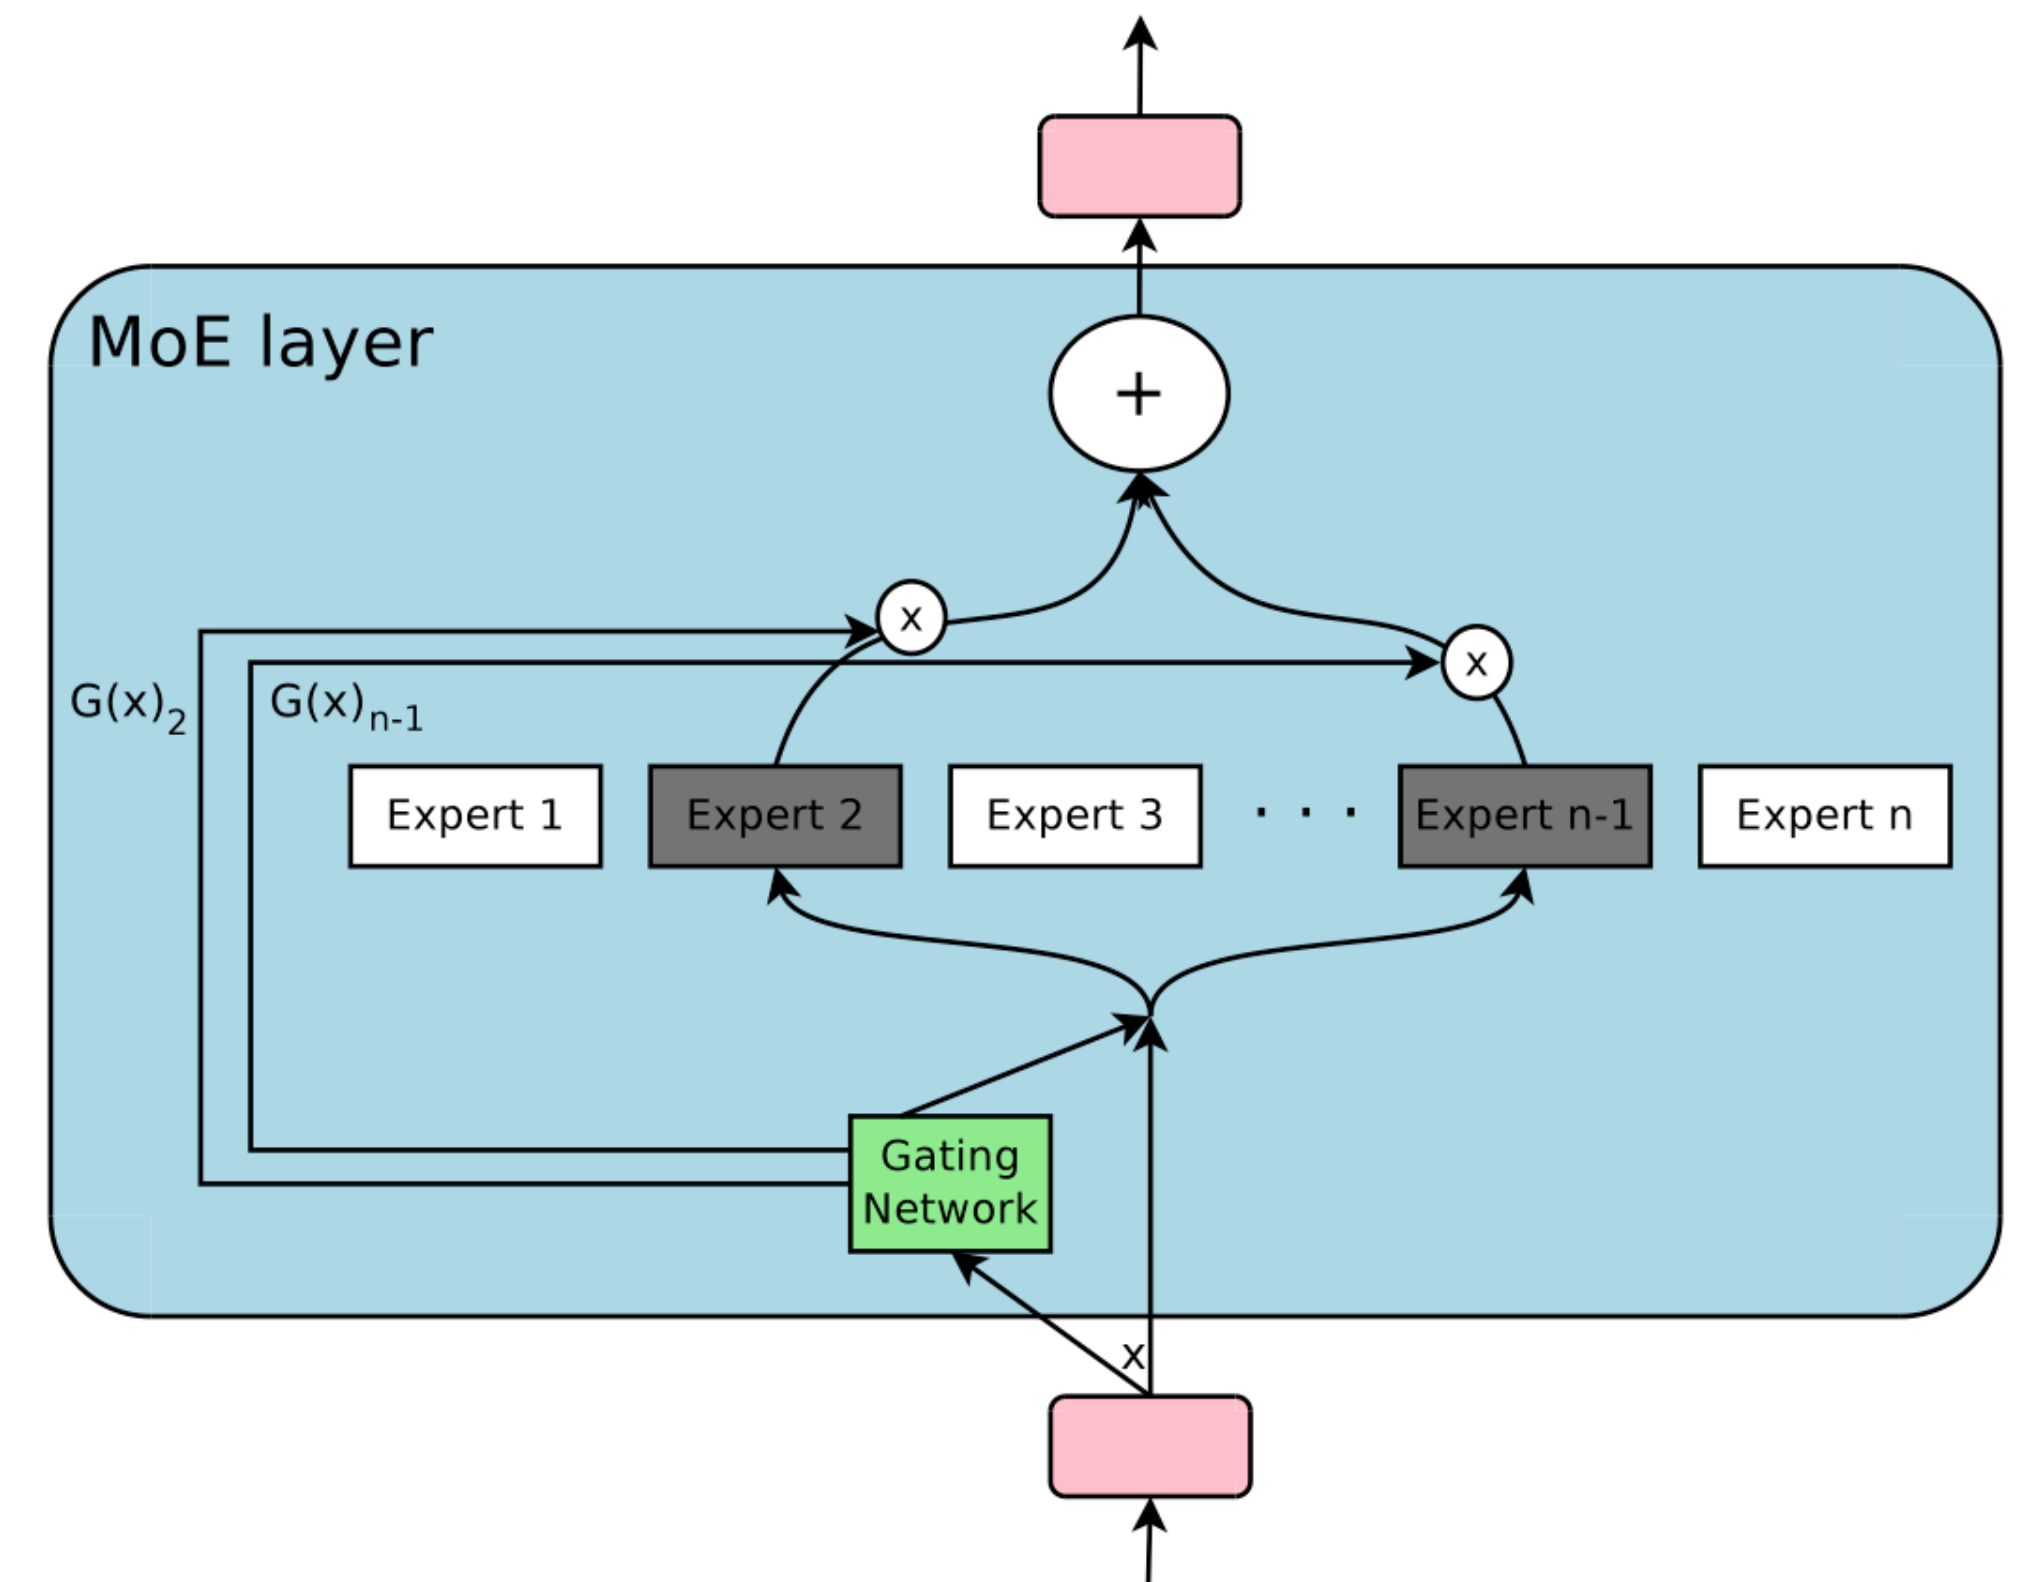
$$\begin{bmatrix} w_1 & 0 & w_3 \\ 0 & w_5 & w_6 \end{bmatrix} \begin{bmatrix} x_1 & x_2 \\ x_3 & x_4 \\ x_5 & x_6 \end{bmatrix}$$

Sparse Weight



$$\begin{bmatrix} w_1 & w_2 & w_3 \\ w_4 & w_5 & w_6 \end{bmatrix} \begin{bmatrix} x_1 & 0 \\ x_3 & x_4 \\ 0 & x_6 \end{bmatrix}$$

Sparse Activation

- Sometimes we sparsify the **activations** instead of weight tensors.
(Often, there are naturally arising sparsity)
- We modify activation functions (FATReLU) or regularizers (Hoyer regularization) to encourage sparsity
- Sometimes we force the activation sparsity, by top-k operations or external routers.
 - The latter variant, called Mixture-of-Experts, are known to be training-efficient.

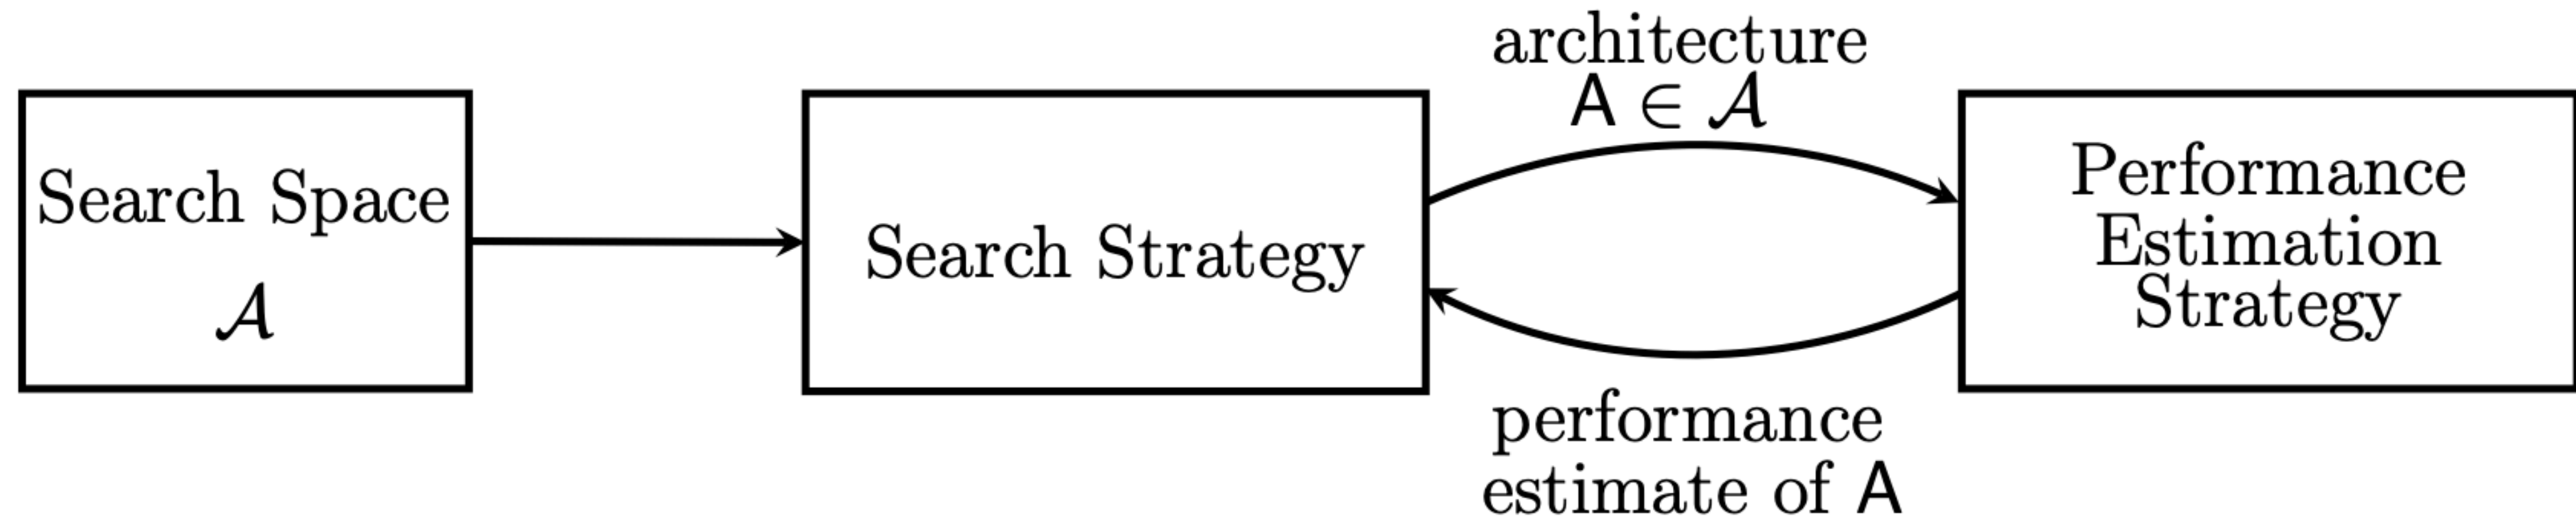


9. Quantization



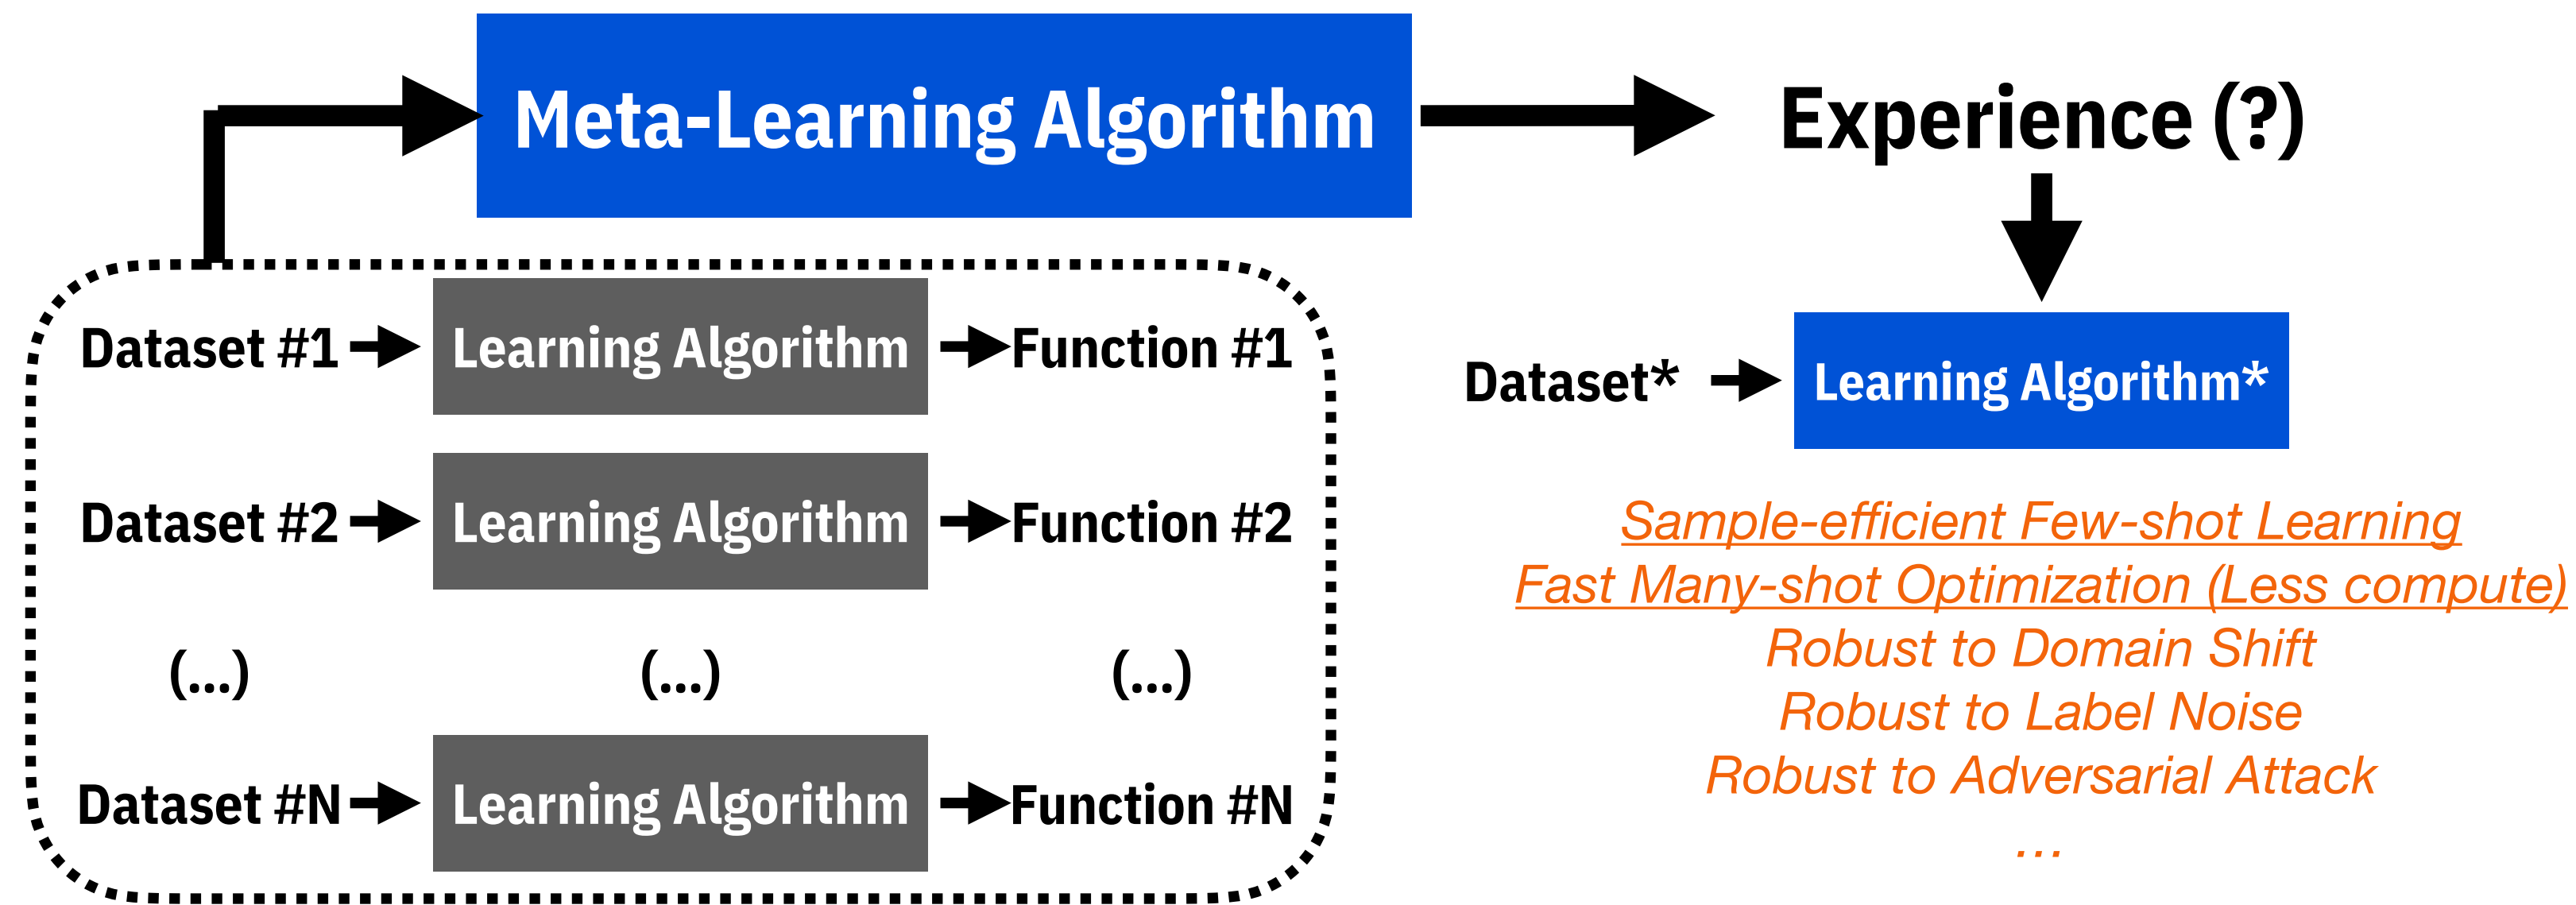
- Dropping the precision of weights/activations to lower bits (e.g., 4/8 bits) for faster compute & smaller size (As we use floating points, this gives rise to many different formats by #bits for exponent, e.g., bfloat16)
- A popular way is “**linear quantization**,” which quantizes weights using uniform grids—key parameters are the scaling and shifting factors.
 - Post-Training Quantization: No special training before/after quantization
 - Quantization-aware training: Additional training, with simulating the noise from quantization.

10. Neural Architecture Search



- Look for the best-performing neural network architecture, within some search space
- Key elements are: Search Space, Search Strategy, and Performance estimation.
 - **Search space:** Networks as a repetition of blocks, where blocks are composed of repeated layers. For each layer, search one specific configuration of efficient modules.
 - **Search strategy:** The decisions are discrete—use RL / evolutionary strategy.
 - **Performance estimation:** Full training is difficult—shorter training or share weights (or zero-cost proxy)

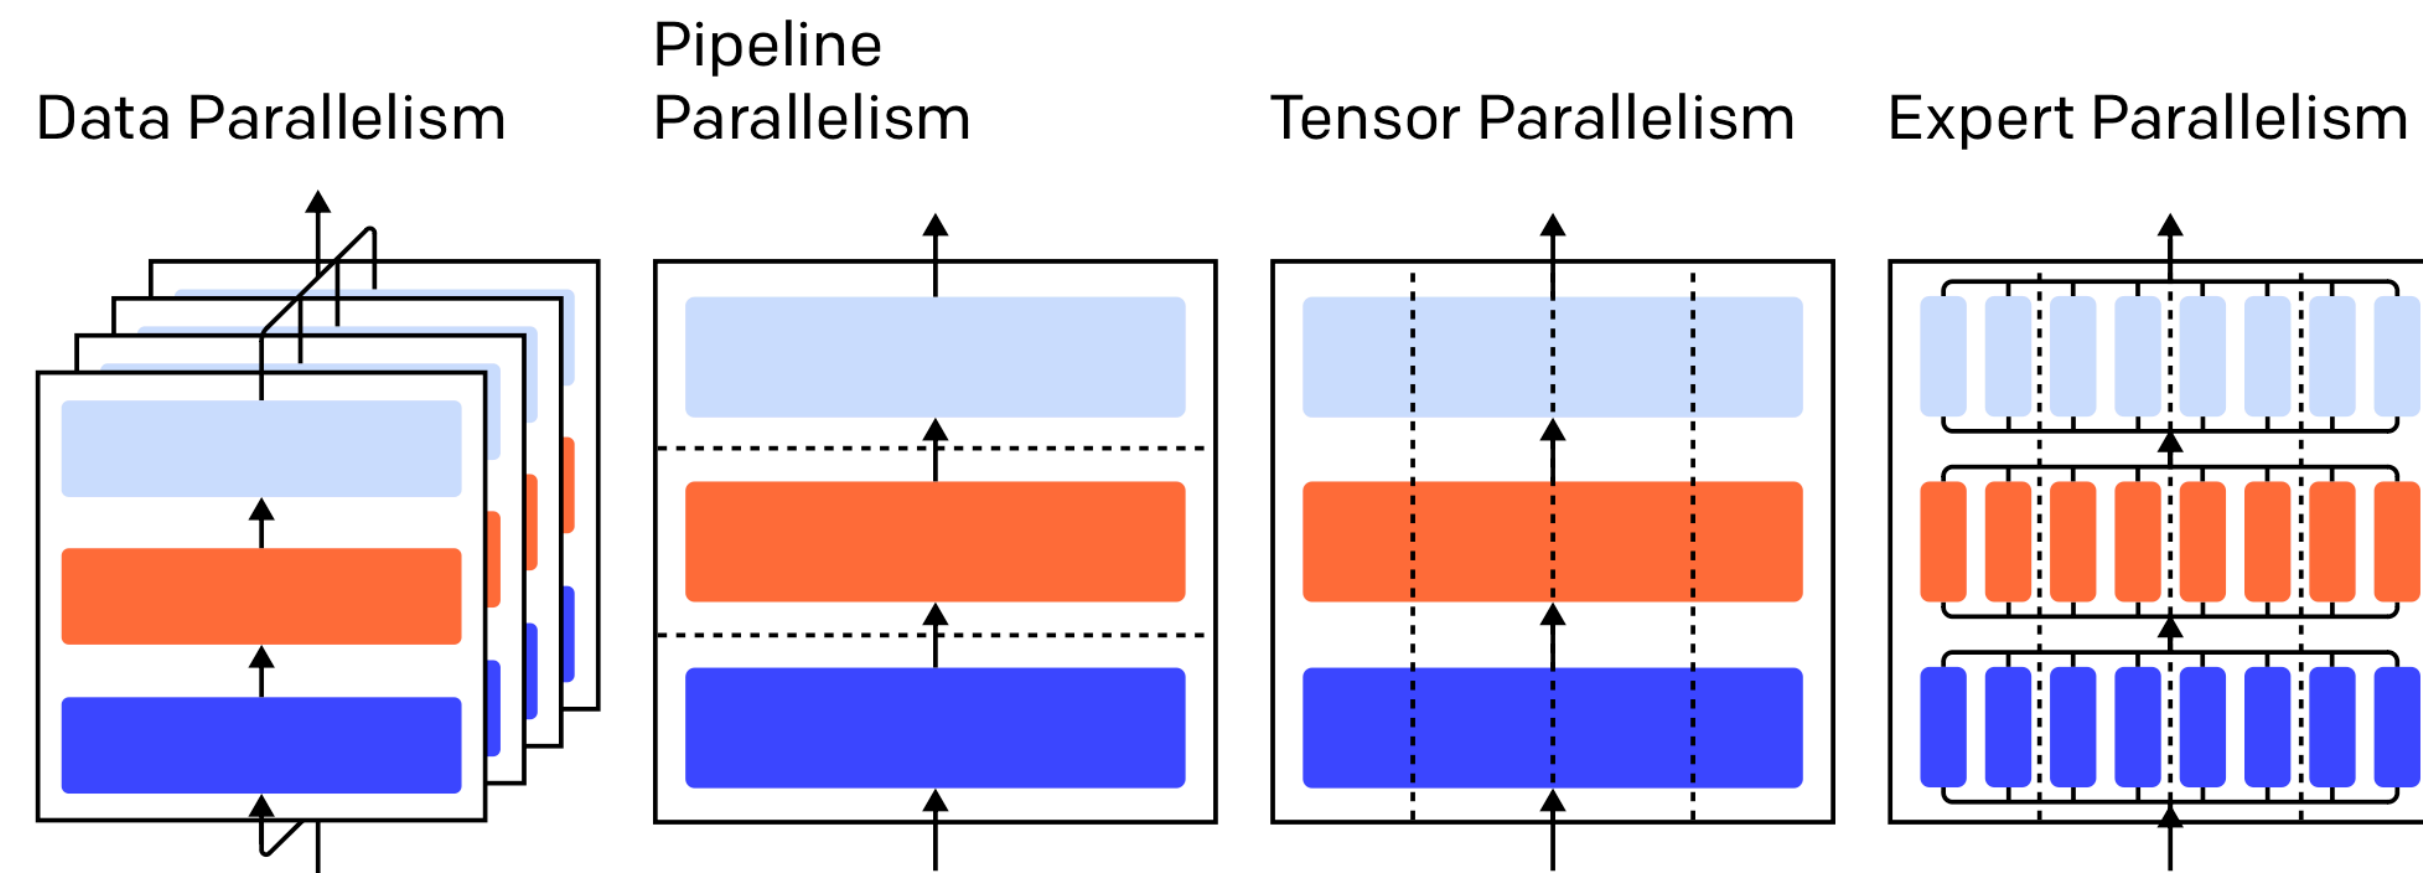
11. Meta-Learning




- For the training-efficiency, transfer the knowledge from task 1, ..., task N to task N+1.
- A popular form is **MAML**—we transfer the weights of a model to a new task, where we train the weights so that they can adapt to a new task within a small number of steps.
- Another interesting application is training the “optimizer” (like Adam).

12. Distributed Training

- For training large-scale models, we exploit many parallelisms.
- Data parallelism:** Data are distributed among workers, Model parameters are shared, Training synchronized by communication.
- Model parallelism:** Model is distributed—each GPU gets
 - (1) some layers (pipeline)
 - (2) some part of each layer (tensor)
 - (3) an expert of MoE (expert)


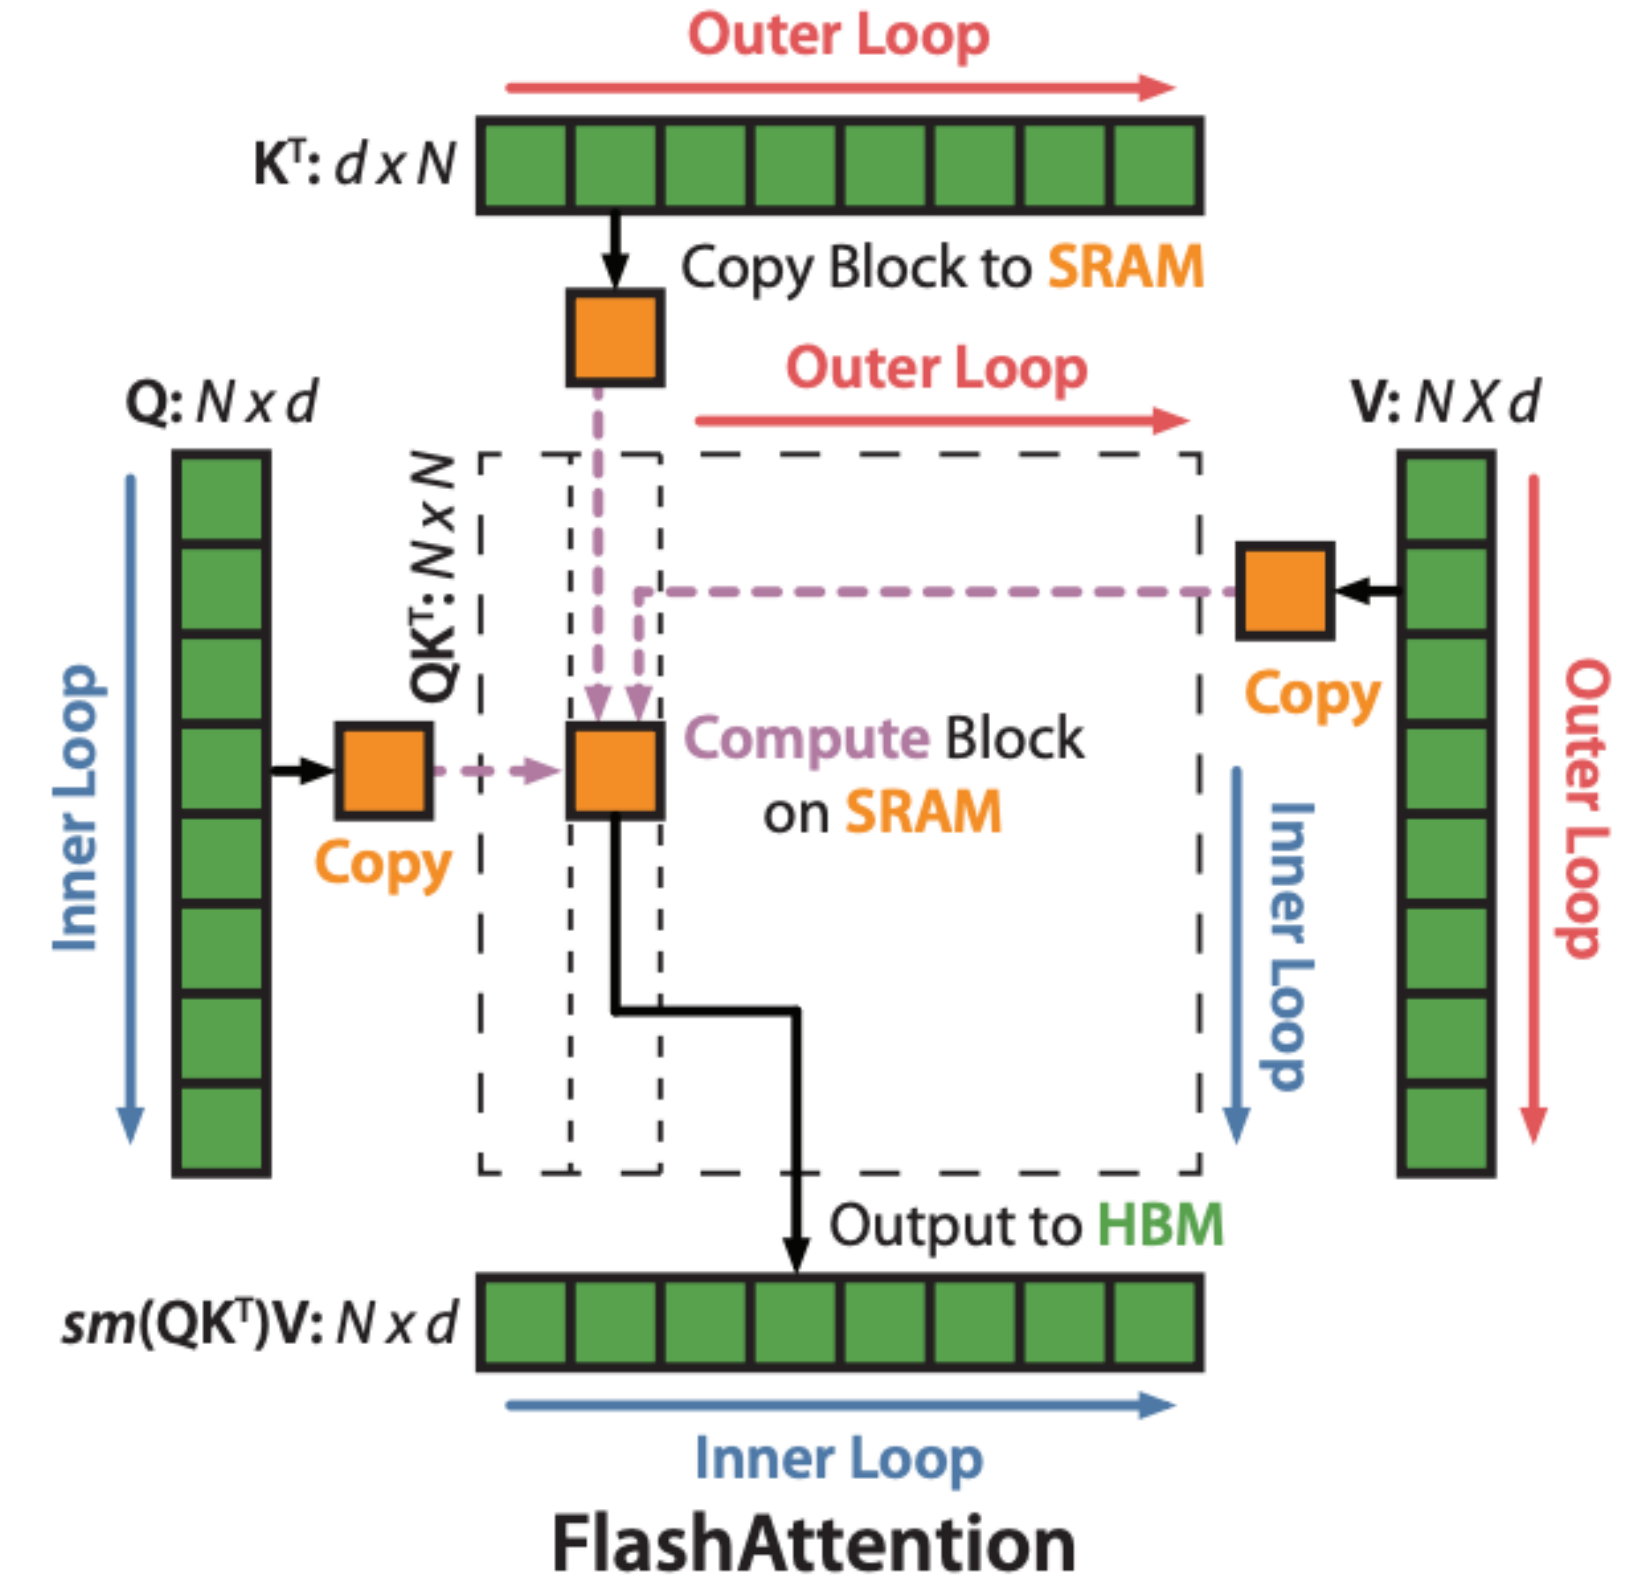


Parameter server (Master)
Coordinates all training




13. Efficient Transformers


- Transformer architectures view any data as a sequence—Compute grows quadratically w.r.t. sequence length!
- Efficient attention** mechanisms use sparse attention among each token inside a sequence
- Flashattention** focuses on optimizing the data movement
- Token pruning/merging** reduce the number of intermediate tokens for a faster inference and training.




(a) Full n^2 attention



(b) Sliding window attention



(c) Dilated sliding window



(d) Global+sliding window