

EECE695D: Efficient ML Systems



Quantization

(Note: Many figures from Song Han's slides)

Quantization

Roughly speaking, quantization is a mapping:

- From an input that belongs to a large set (e.g., \mathbb{R} , 32-bit floating point, ...)
- To an output that belongs to a smaller, discrete set (e.g., FP32, 8-bit integer, ...)



Quantization in Deep Learning

Typically a scalar quantization, dropping the bitwidth of the weights/activations (e.g., 32 bits -> 16/8/4 bits).
Enjoys many benefits, including:

- Multiplying low-precision weights with low-precision activation requires **less compute**
- Loading low-precision weights/activations requires **less memory bandwidth usage**
- Storing quantized data/models requires **less storage space**
- Low bit processing typically requires **less silicon space!**

Add energy (pJ)		Mem access energy (pJ)		Add area (μm ²)	
INT8	FP32	Cache (64-bit)		INT8	FP32
0.03	0.9	8KB		36	4184
30X energy reduction		32KB		116X area reduction	
		1MB			
Mult energy (pJ)		DRAM		Mult area (μm ²)	
INT8	FP32	1300-2600		INT8	FP32
0.2	3.7	Up to 4X energy reduction		282	7700
18.5X energy reduction				27X area reduction	

TSMC45nm, 0.9V

Recap: Data numerics

Integers. No fractions, just integers.

- *Unsigned Integer.* Can represent $\{0, 1, \dots, 2^n - 1\}$
- *Signed Integer (ver 1. Sign-magnitude).*
 - Can represent $\{-2^{n-1} - 1, \dots, 2^{n-1} - 1\}$
 - Both $000\dots00$ and $100\dots00$ represents 0
- *Signed Integer (ver 2. 2's complement).*
 - Can represent $\{-2^{n-1}, \dots, 2^{n-1} - 1\}$
 - $100\dots00$ now represents -2^{n-1} instead of 0

0	0	1	1	0	0	0	1
x	x	x	x	x	x	x	x

$$2^7 + 2^6 + 2^5 + 2^4 + 2^3 + 2^2 + 2^1 + 2^0 = 49$$

Sign Bit

1	0	1	1	0	0	0	1
	x	x	x	x	x	x	x

$$- 2^6 + 2^5 + 2^4 + 2^3 + 2^2 + 2^1 + 2^0 = -49$$

1	1	0	0	1	1	1	1
x	x	x	x	x	x	x	x

$$-2^7 + 2^6 + 2^5 + 2^4 + 2^3 + 2^2 + 2^1 + 2^0 = -49$$

Fixed-Point. Integers, but with fractions.

- Bits after the virtual decimal point represent fractions
- Mostly used under very specific circumstances (e.g., super low-cost microprocessors)
Used in very early works (e.g., Vanhoucke et al., 2011; Hwang & Sung, 2014)



Integer . Fraction

“Decimal” Point



x x x x x x x x

$$-2^3 + 2^2 + 2^1 + 2^0 + 2^{-1} + 2^{-2} + 2^{-3} + 2^{-4} = 3.0625$$




x x x x x x x x

$$(-2^7 + 2^6 + 2^5 + 2^4 + 2^3 + 2^2 + 2^1 + 2^0) \times 2^{-4} = 49 \times 0.0625 = 3.0625$$

Floating-Point. Fixed-point, but the scaling factor is not fixed

- Consists of a **sign bit**, **exponent bits**, and **fraction bits** (a.k.a. mantissa)
- Exponent = Dynamic range (important for accumulation), Fraction = Precision




Sign 8 bit Exponent

23 bit Fraction


$$(-1)^{\text{sign}} \times (1 + \mathbf{\text{Fraction}}) \times 2^{\text{Exponent}-127} \quad \leftarrow \quad \text{Exponent Bias} = 127 = 2^{8-1}-1$$

(significant / mantissa)

$$0.265625 = 1.0625 \times 2^{-2} = (1 + \underline{0.0625}) \times 2^{\underline{125}-127}$$



IEEE 754 Single Precision 32-bit Float (IEEE FP32)



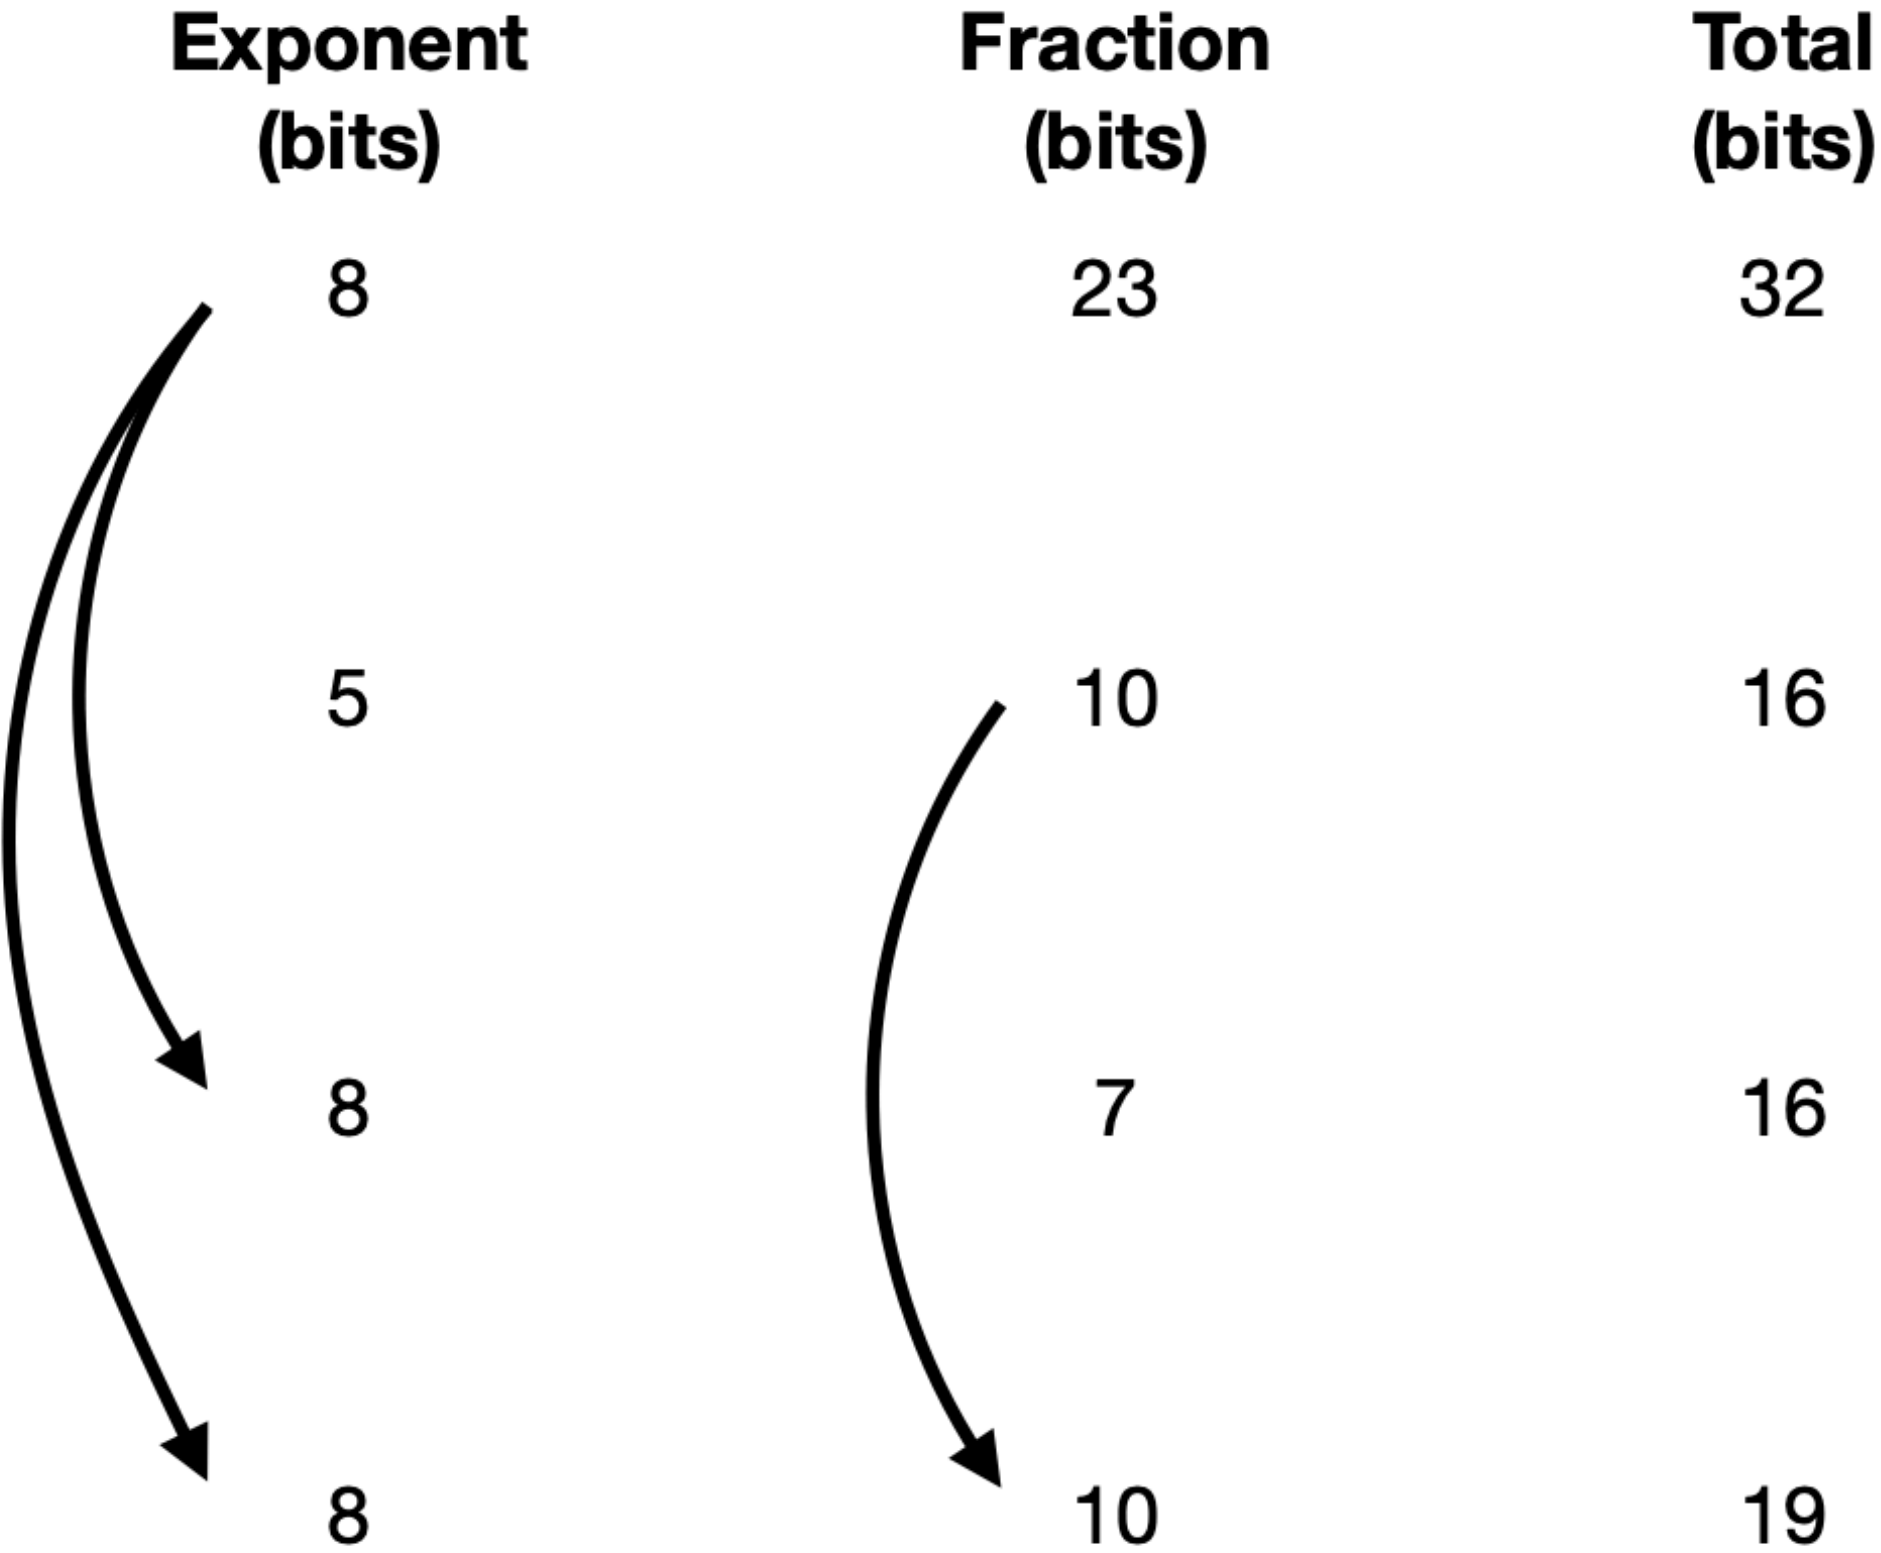

IEEE Half Precision 16-bit Float (IEEE FP16)



Brain Float (BF16)



Nvidia TensorFloat (TF32)



BF16. Can replace/be-combined-with FP32 (identical underflow, overflow, NaN, ...)

Typical MAC: Multiply in bfloat16, Accumulate in FP32

Usually no “loss scaling” required

Smaller mantissa than FP16 -> smaller size needed in silicon

TF32. Exponent of FP32, Mantissa of FP16.

Used in NVIDIA Ampere.

More fraction -> better precision -> better performance.

Naturally, slower than BF/FP16

FP32	TF32	FP16 / BF16
1x	8x	16x

Table 1. Relative throughput of A100 GPU math.

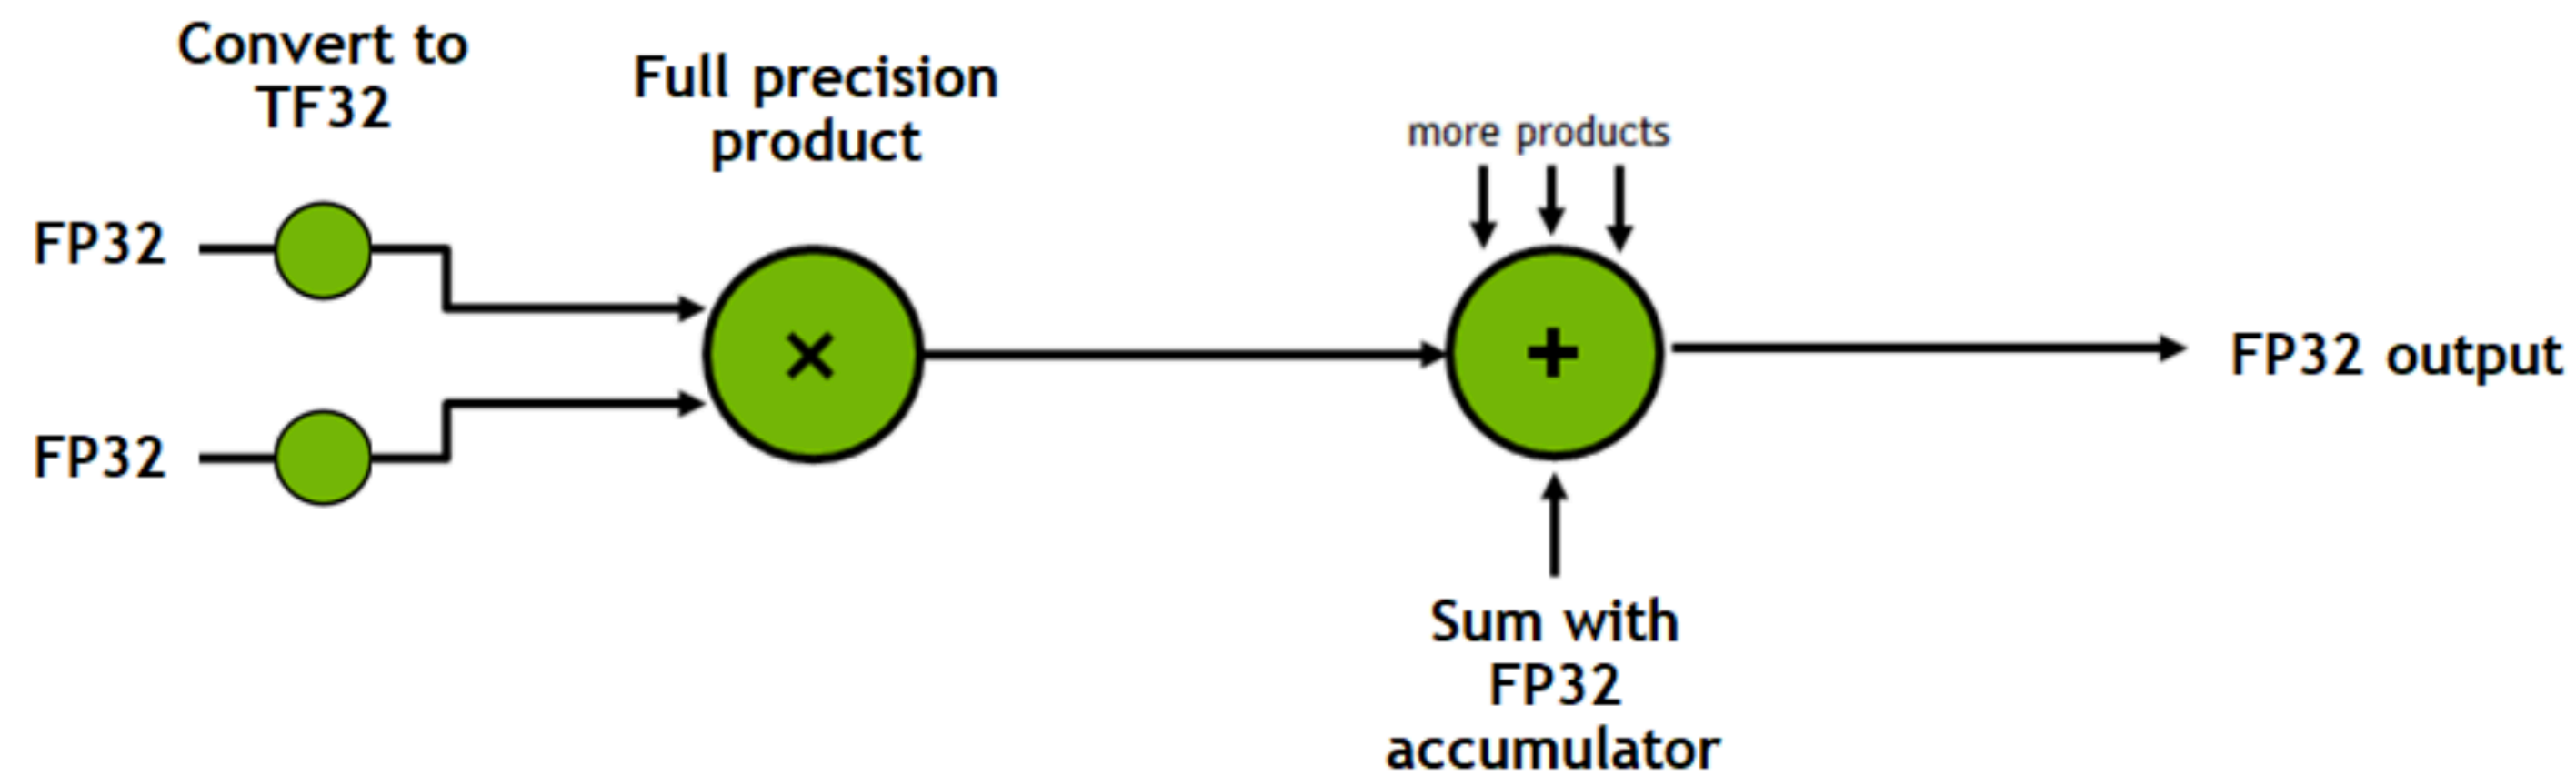



Figure 1. Ampere A100 Tensor Core operation.




K-means based quantization

Around 2016, fancy methods based on K-means gained popularity (e.g., Han et al. 2016)


(1D) K-Means. Given $x_1, \dots, x_n \in \mathbb{R}$, find $c_1, \dots, c_k \in \mathbb{R}$ that minimizes

$$\frac{1}{n} \sum_{i=1}^n \min_{j \in [k]} (x_i - c_j)^2$$

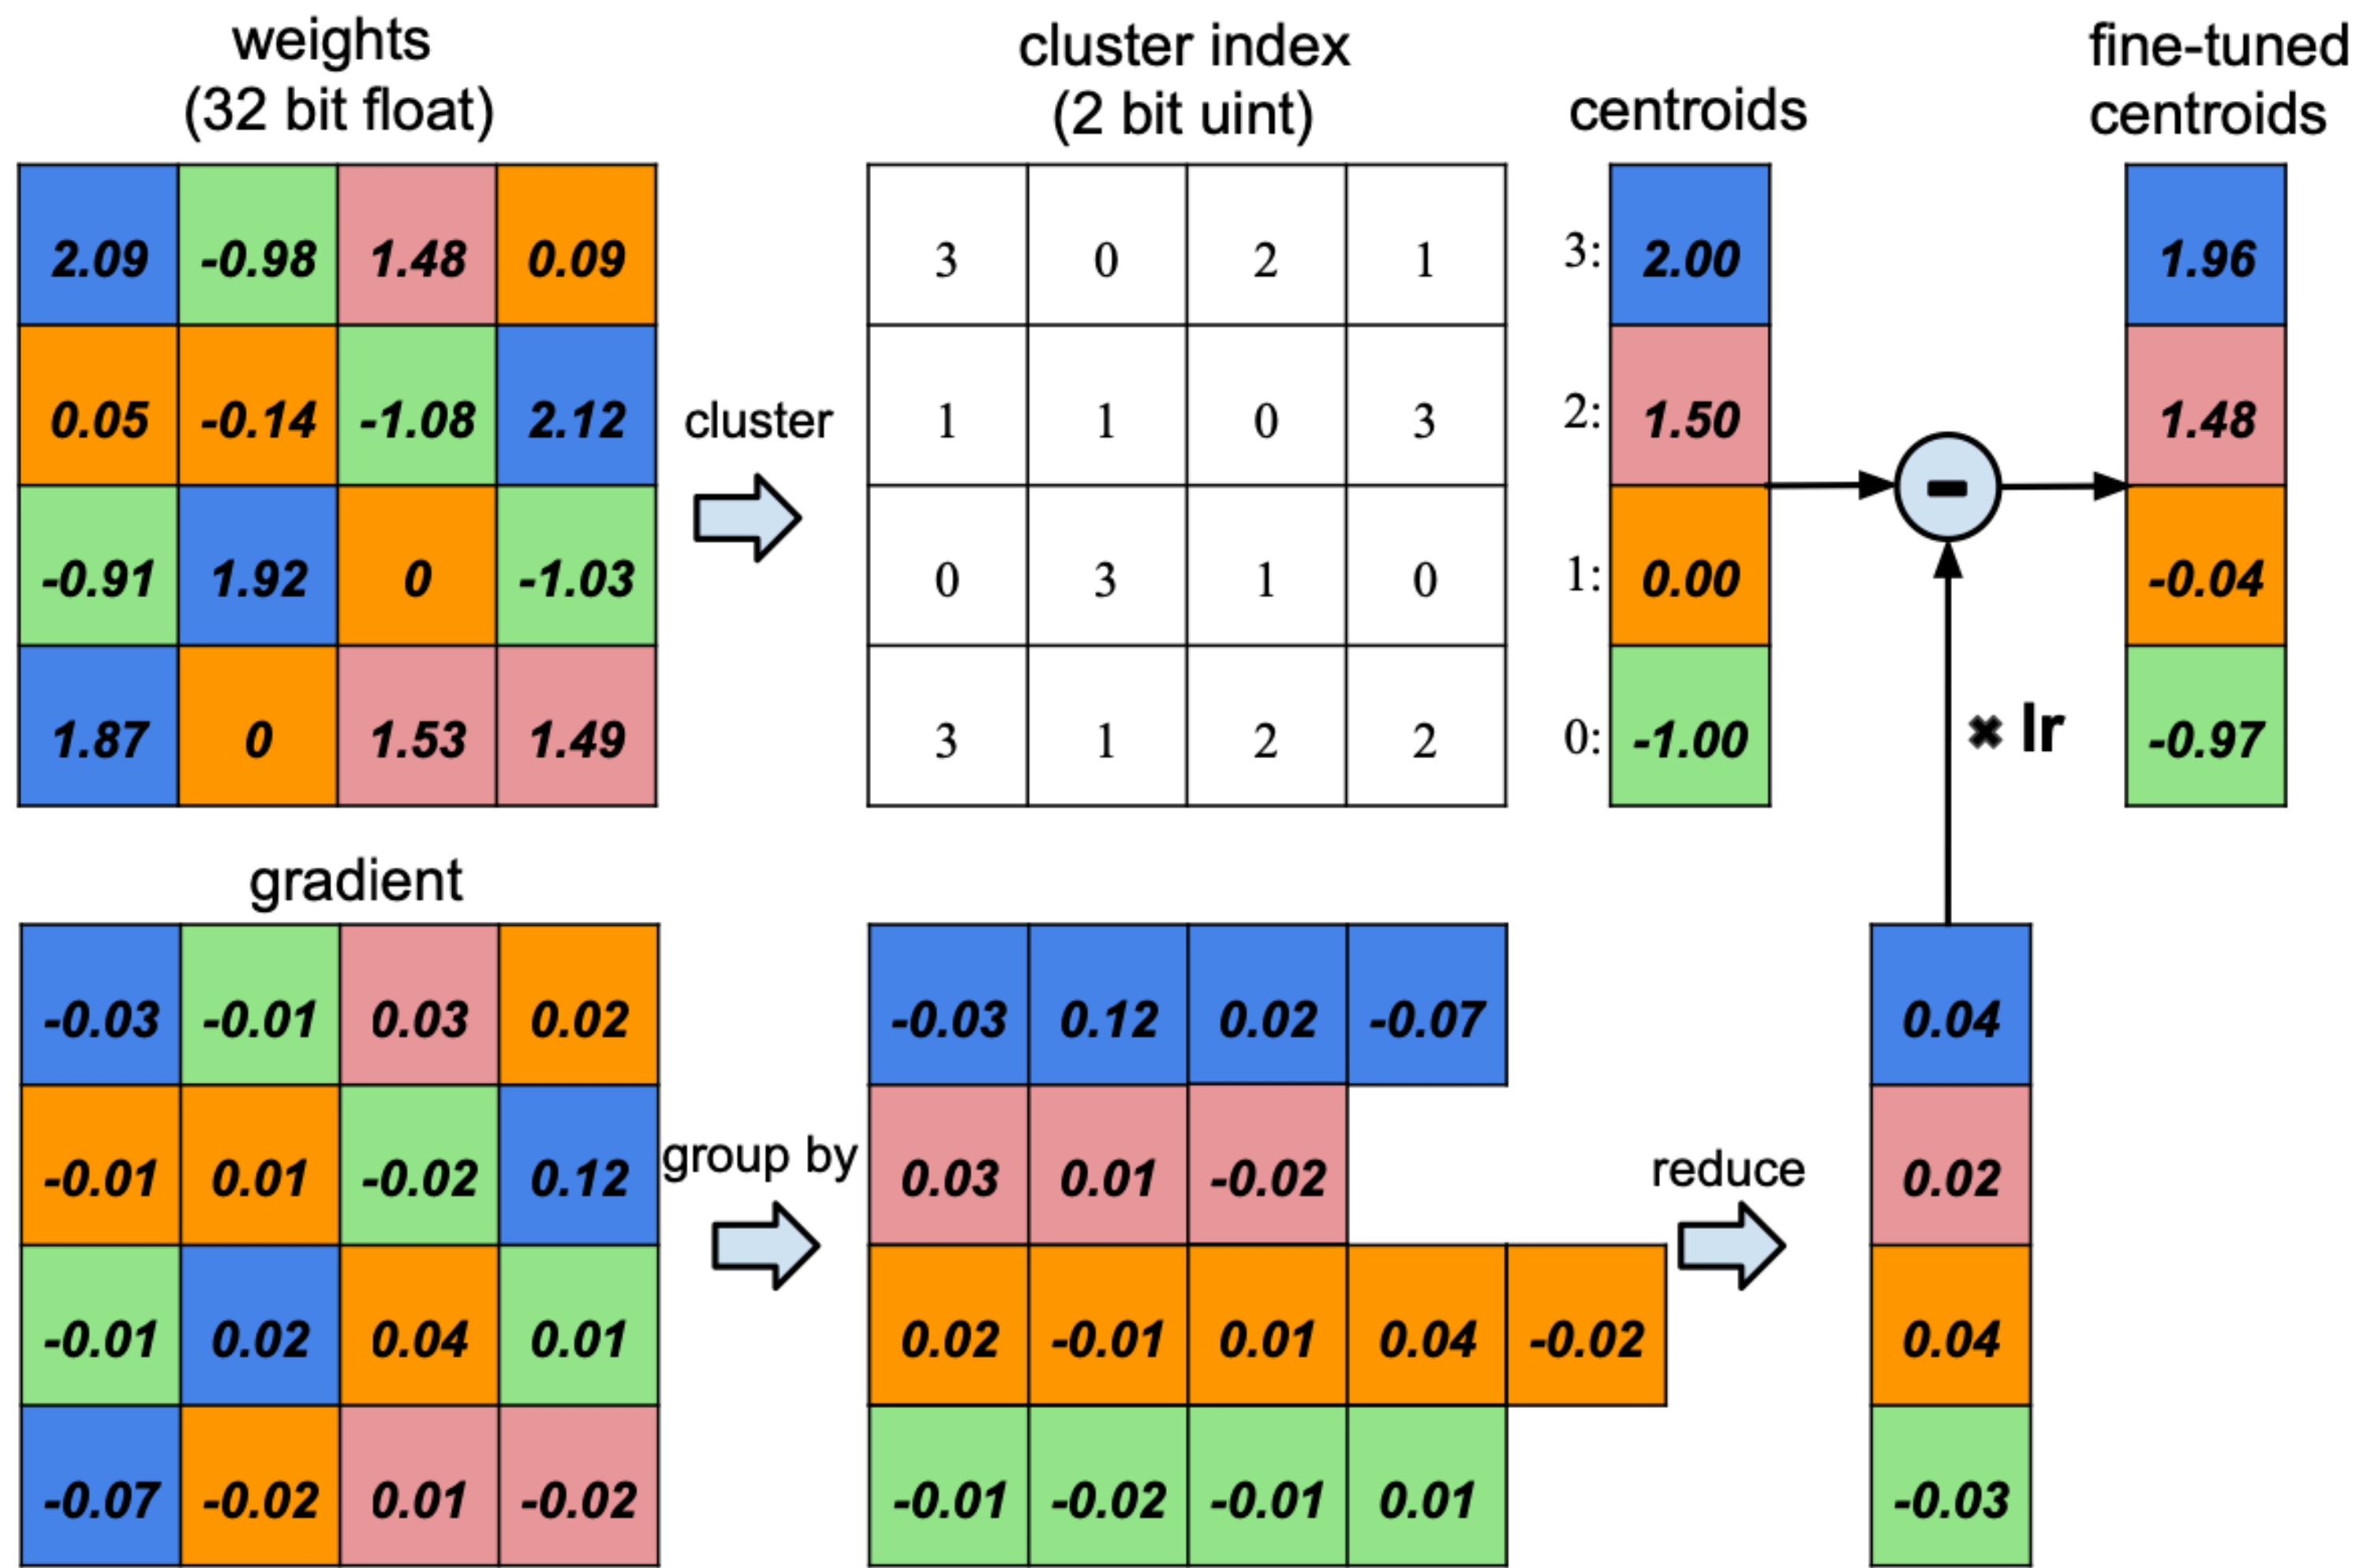
Famous methods to solve this include: Lloyd's algorithm, k-means++, ...



Idea. Do the same thing for the weight matrix of a neural net
Need to store both index and codebook.




Codebook training. Gradient updates can be done as in convolutional neural networks.
Note: Not for training neural net from scratch; only for compressing model itself!



Limitation. Reduces weight storage (INT index, FP codebook) but not compute/activation (uses FP arithmetics).
Usually a common drawback of “nonlinear” methods.

Linear quantization

More mainstream nowadays; weight matrix \mathbf{r} is represented as $\mathbf{r} = S(\mathbf{q} - \mathbf{Z})$ (with S, Z being scalar)
Typically introduces more quantization error than nonlinear (but benefits usually outweigh)



Floating-point

Integer

Integer

Floating-point

- quantization parameter
- allow real number $r=0$ be exactly representable by a quantized integer Z
- quantization parameter

Note. Having an exact zero is important!
(we've seen naturally arising sparsity)

Quant vs. Dequant. The previous formula

$$\mathbf{r} = S(\mathbf{q} - Z)$$

is actually what we call dequantization.

The formula for quantization can be written as:

$$\mathbf{q} = \text{Clip} \left(\text{Round} \left(\frac{\mathbf{r}}{S} \right) + Z \right)$$

Q. Does how-to-round matter? Can we improve by loss-oriented rounding?



A. Yes, see AdaRound paper by Nagel et al. (2020)

Rounding scheme	Acc(%)
Nearest	52.29
Ceil	0.10
Floor	0.10
Stochastic	52.06±5.52
Stochastic (best)	63.06

Table 1. Comparison of ImageNet validation accuracy among different rounding schemes for 4-bit quantization of the first layer of Resnet18. We report the mean and the standard deviation of 100 stochastic (Gupta et al., 2015) rounding choices (Stochastic) as well as the best validation performance among these samples (Stochastic (best)).

- Q.** Given a matrix **r**, How should we optimize the scale *S* and zero *Z*?
- A.** Typical method: Decide *S* to match max & min, then decide *Z* to be the nearest point.


2.09	-0.98	1.48	0.09
0.05	-0.14	-1.08	2.12
-0.91	1.92	0	-1.03
1.87	0	1.53	1.49



Binary	Decimal
01	1
00	0
11	-1
10	-2

$$\begin{aligned} S &= \frac{r_{\max} - r_{\min}}{q_{\max} - q_{\min}} \\ &= \frac{2.12 - (-1.08)}{1 - (-2)} \\ &= 1.07 \end{aligned}$$

- Q.** Given a matrix \mathbf{r} , How should we optimize the scale S and zero Z ?
- A.** Typical method: Decide S to match max & min, then decide Z to be the nearest point.



2.09	-0.98	1.48	0.09
0.05	-0.14	-1.08	2.12
-0.91	1.92	0	-1.03
1.87	0	1.53	1.49

$$\begin{aligned}
 Z &= q_{\min} - \frac{r_{\min}}{S} \\
 &= \text{round}\left(-2 - \frac{-1.08}{1.07}\right) \\
 &= -1
 \end{aligned}$$

Binary	Decimal
01	1
00	0
11	-1
10	-2

Dot products can now be done more efficiently.

Suppose that we are doing

$$Y = W^{\top}X, \quad W, X \in \mathbb{R}^d, Y \in \mathbb{R}^1$$

where we have $W = S_W(\mathbf{q}_W - Z_W)$, $X = S_X(\mathbf{q}_X - Z_X)$, and $Y = S_Y(\mathbf{q}_Y - Z_Y)$.

(Here, let Z be a vectorized form $Z = z \cdot \mathbf{1}$)

In other words, we are doing

$$S_Y(\mathbf{q}_Y - Z_Y) = (S_W S_X) \cdot (\mathbf{q}_W \mathbf{q}_X - Z_X \mathbf{q}_W - Z_W \mathbf{q}_X + Z_W Z_X)$$

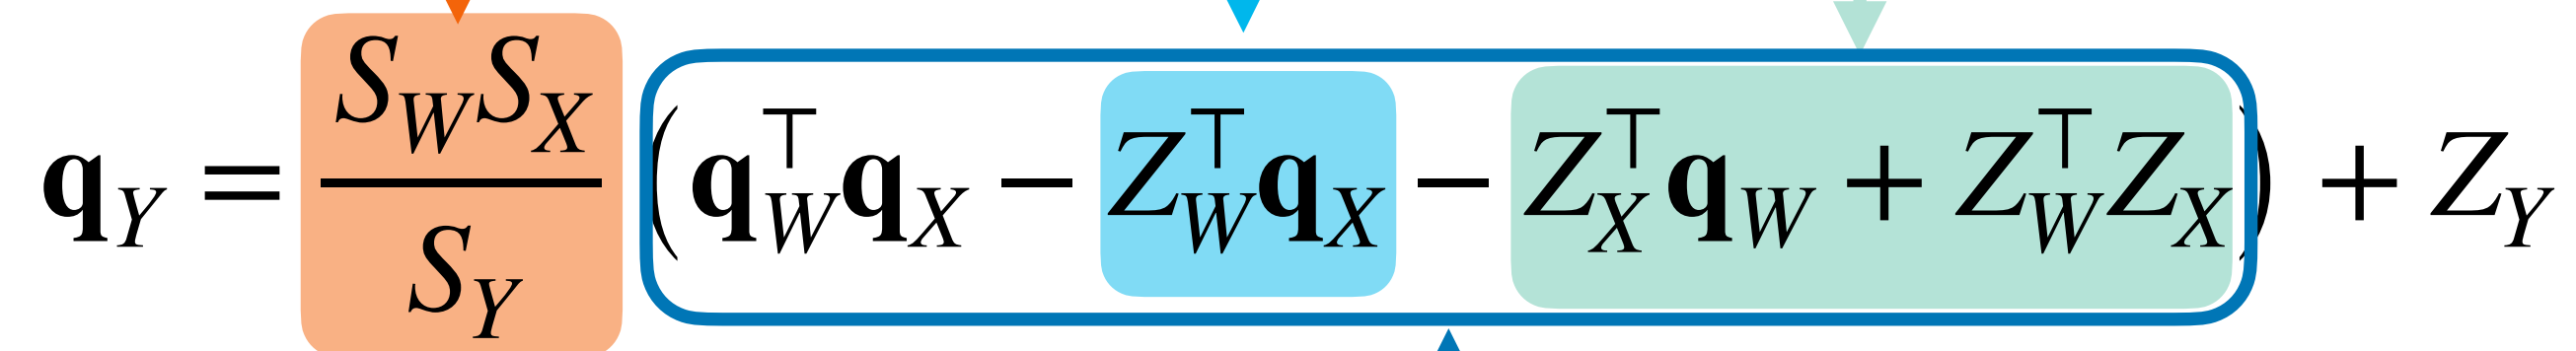
Assume that we are not dynamically adjusting S, Z . Then, \mathbf{q}_Y can be written as:

$$\mathbf{q}_Y = \frac{S_W S_X}{S_Y} (\mathbf{q}_W^{\top} \mathbf{q}_X - Z_W^{\top} \mathbf{q}_X - Z_X^{\top} \mathbf{q}_W + Z_W^{\top} Z_X) + Z_Y$$

Much easier when
 $Z_W = 0$ (symmetric quantization!)

Should be rescaled to
a low-bit integer
value empirically in (0,1)

Can be pre-computed
and folded into bias




The diagram shows the equation $\mathbf{q}_Y = \frac{S_W S_X}{S_Y} (\mathbf{q}_W^\top \mathbf{q}_X - Z_W^\top \mathbf{q}_X - Z_X^\top \mathbf{q}_W + Z_W^\top Z_X) + Z_Y$. Annotations include: an orange box around $\frac{S_W S_X}{S_Y}$ with an arrow from 'Should be rescaled to a low-bit integer value empirically in (0,1)'; a blue box around $Z_W^\top \mathbf{q}_X$ with an arrow from 'Much easier when $Z_W = 0$ (symmetric quantization!)'; a green box around $Z_X^\top \mathbf{q}_W + Z_W^\top Z_X$ with an arrow from 'Can be pre-computed and folded into bias'; and a blue box around the entire parentheses with an arrow from 'Multiplications done in low-bit integer Accumulations done in 32-bit integer'.

$$\mathbf{q}_Y = \frac{S_W S_X}{S_Y} (\mathbf{q}_W^\top \mathbf{q}_X - Z_W^\top \mathbf{q}_X - Z_X^\top \mathbf{q}_W + Z_W^\top Z_X) + Z_Y$$


Multiplications done in low-bit integer
Accumulations done in 32-bit integer

Symmetric Q. Uses the range $[-|r|_{\max}, +|r|_{\max}]$ instead of $[r_{\min}, r_{\max}]$



Note. For activations, using ReLU means only nonnegatives will be there, i.e., wasted range.


Full-Range. Truncate the max



$$S = \frac{|r|_{\max}}{2^{N-1}}$$

(ONNX, PyTorch, ...)

Restricted-Range. Drop a symbol



$$S = \frac{|r|_{\max}}{2^{N-1} - 1}$$

(TensorFlow, NVIDIA TensorRT, ...)

